

**Digital Design and
Computer Architecture LU**

**Digital Design and Computer Architecture
MiMi – Minimal MIPS
Exercise III**

{fhuemer, fkriebel, jmaier}@ecs.tuwien.ac.at
Department of Computer Engineering
University of Technology Vienna

Vienna, April 30, 2018

1 Introduction

This document describes MiMi, a minimal MIPS implementation. It is mostly binary compatible with the original MIPS implementation, as described in the seminal *Computer Architecture: A Quantitative Approach* [1] and *Computer Organization and Design* [2]. However, only a subset of the instruction set is implemented. Also, the design is a Harvard architecture, which entails that exception and interrupt handling is slightly different than in the original MIPS implementation.

Figure 1.1 shows the 5-stage pipeline of the processor to be implemented. It comprises five stages: fetch, decode, execute, memory and write-back. The data path is drawn in black; signals that flush a pipeline stage are blue, signals that stall the pipeline are green, and signals that refer to exceptions are red. In the upcoming assignments, the parts to be implemented will be shown in light blue and entities to be instantiated will be shaded, to ease your navigation through the design.




Figure 1.1: MiMi pipeline

Section 2 describes the implementation of the basic elements of the pipeline, such as the ALU and the register file. These elements are put together to form a pipeline in Section 3.

The questions in the **Theory** sections do not need to be answered in order to achieve points for the practical part. However, they may be among the theoretical questions in the final exam.

Contents

1	Introduction	1
2	Level 0: Basic Elements	4
2.1	ALU	5
2.2	Jump Unit	7
2.3	Memory Unit	8
2.4	Register File	12
3	Level 1: Pipeline	13
3.1	Fetch	14
3.2	Decode	15
3.3	Execute	20
3.4	Memory	22
3.5	Write-Back	24
3.6	Pipeline	25
A	Tools	26
B	Automated Test Environment	27
C	Submission Requirements	28
C.1	Exercise III	28

List of Tables

2.1	ALU interface	6
2.2	ALU result computation	6
2.3	ALU zero-flag computation	6
2.4	ALU overflow conditions	6
2.5	Jump Unit interface	7
2.6	Jump Unit operations	7
2.7	Memory Unit interface	9
2.8	MEM_OP_TYPE fields	9
2.9	MEM_OUT_TYPE fields	9
2.10	Computation of M.byteena and M.wrdata, W = DCBA	9
2.11	Computation of R, D = DCBA	10
2.12	Memory load exception computation	10
2.13	Memory store exception computation	11
2.14	Register file interface	12
3.1	Fetch stage interface	14
3.2	Decode stage interface	16
3.3	EXEC_OP_TYPE fields	16
3.4	COP0_OP_TYPE fields	16
3.5	WB_OP_TYPE fields	17
3.6	MiMi instructions	18
3.7	MiMi special instructions	19
3.8	MiMi regimm instructions	19
3.9	MiMi cop0 instructions	19
3.10	Execute stage interface	21
3.11	Memory stage interface	23
3.12	Write-back stage interface	24
3.13	Pipeline interface	25

List of Figures

1.1	MiMi pipeline	1
3.1	Instruction formats	17

Listings

1	Assembler example without forwarding	25
2	Makefile fragment	26

2 Level 0: Basic Elements

This assignment consists of four relatively simple hardware units. Implement the units described in this section, and write appropriate test benches. Test the units thoroughly, as errors introduced at this stage might be very difficult to find in later stages.

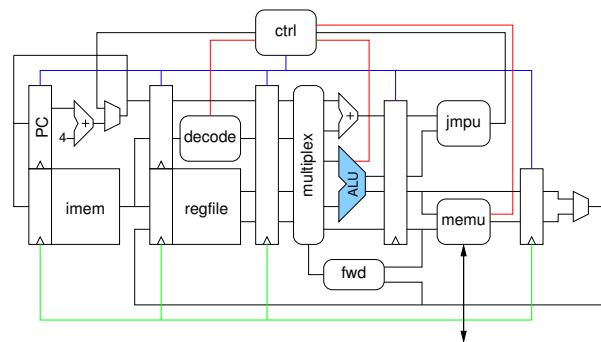
Points 5

Evaluation

The assignment will be evaluated with test benches, which thoroughly test the individual components. Points will be granted if the test benches are passed successfully. The assignment is part of lab exercise III.

2.1 ALU

alu.vhd



Description

The arithmetic logic unit (ALU) carries out – as its name suggests – arithmetic and logic operations. The interface of the ALU is described by Table 2.1; it shall implement the operations described in Table 2.2. The computation of the zero flag Z and the overflow flag V is shown in Tables 2.3 and 2.4, respectively. Note that the shift operations can be implemented conveniently with the functions `shift_left()` and `shift_right` from the package `numeric_std`.

Theory

A general comparison between two values for less-than/greater-than requires a subtraction and an appropriate evaluation of the most significant bits of the result. How many logic elements does the critical path for a comparison of two n -bit values contain asymptotically (i.e., $O(n^2)$, $O(n)$, $O(\log n)$, ...)? What about comparing for equality/inequality? Why is a comparison for less-than-zero cheap when using a two's complement representation?

Signal	Direction	Type	Width	Description
op	in	ALU_OP_TYPE	–	Operation
A	in	std_logic_vector	DATA_WIDTH	Operand A
B	in	std_logic_vector	DATA_WIDTH	Operand B
R	out	std_logic_vector	DATA_WIDTH	Result
Z	out	std_logic	–	Zero flag
V	out	std_logic	–	Overflow flag

Table 2.1: ALU interface

op	R
ALU_NOP	A
ALU_LUI	B sll 16
ALU_SLT	A < B ? 1 : 0, signed
ALU_SLTU	A < B ? 1 : 0, unsigned
ALU_SLL	B sll A(DATA_WIDTH_BITS-1 downto 0)
ALU_SRL	B srl A(DATA_WIDTH_BITS-1 downto 0)
ALU_SRA	B sra A(DATA_WIDTH_BITS-1 downto 0)
ALU_ADD	A + B
ALU_SUB	A - B
ALU_AND	A and B
ALU_OR	A or B
ALU_XOR	A xor B
ALU_NOR	not (A or B)

Table 2.2: ALU result computation

op	Z
ALU_SUB	if A = B then Z <= '1'; else Z <= '0'; end if ;
otherwise	if A = 0 then Z <= '1'; else Z <= '0'; end if ;

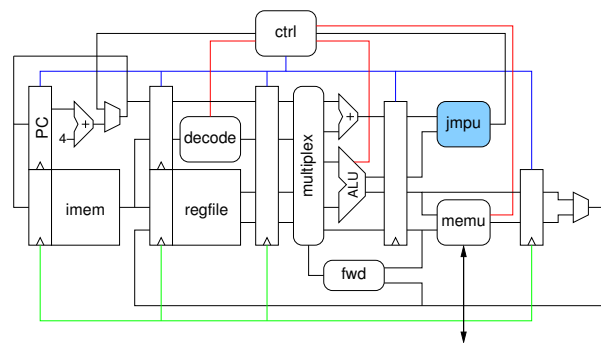
Table 2.3: ALU zero-flag computation

op	A	B	R	V
ALU_ADD	≥ 0	≥ 0	< 0	'1'
ALU_ADD	< 0	< 0	≥ 0	'1'
ALU_SUB	≥ 0	< 0	< 0	'1'
ALU_SUB	< 0	≥ 0	≥ 0	'1'
	otherwise			'0'

Table 2.4: ALU overflow conditions

2.2 Jump Unit

jmpu.vhd



Description

The interface of the jump unit is shown in Table 2.5. It shall implement the operations shown in Table 2.6. The zero flag Z corresponds to the zero flag of the ALU, while the negative flag N corresponds to a negative result from the ALU.

Theory

Table 2.6 does not contain operations for all boolean combinations of N and Z. Would an operation for N **and** Z make sense? If so, what would be the high-level comparison? If not, explain why.

Signal	Direction	Type	Description
op	in	JMP_OP_TYPE	Operation
N	in	std_logic	Negative flag
Z	in	std_logic	Zero flag
J	out	std_logic	Jump

Table 2.5: Jump Unit interface

Operation	J
JMP_NOP	'0'
JMP_JMP	'1'
JMP_BEQ	Z
JMP_BNE	not Z
JMP_BLEZ	N or Z
JMP_BGTZ	not (N or Z)
JMP_BLTZ	N
JMP_BGEZ	not N

Table 2.6: Jump Unit operations

Signal	Direction	Type	Width	Description
op	in	MEM_OP_TYPE	–	Access type
A	in	std_logic_vector	ADDR_WIDTH	Address
W	in	std_logic_vector	DATA_WIDTH	Write data
D	in	std_logic_vector	DATA_WIDTH	Data from memory
M	out	MEM_OUT_TYPE	–	Interface to memory
R	out	std_logic_vector	DATA_WIDTH	Result of memory load
XL	out	std_logic	–	Load exception
XS	out	std_logic	–	Store exception

Table 2.7: Memory Unit interface

Field	Type	Description
memread	std_logic	Read from memory
memwrite	std_logic	Write to memory
memtype	MEMTYPE_TYPE	Word, half-word or byte access

Table 2.8: MEM_OP_TYPE fields

Field	Type	Width	Description
address	std_logic_vector	ADDR_WIDTH	Address to read from or write to
rd	std_logic	–	Asserted for reads
wr	std_logic	–	Asserted for writes
byteena	std_logic_vector	4	Byte-enable signal for sub-word writes
wrdata	std_logic_vector	DATA_WIDTH	Data to be written

Table 2.9: MEM_OUT_TYPE fields

Operation	A(1 downto 0)	M.byteena	M.wrdata
MEM_B MEM_BU	"00"	"1000"	AXXX
	"01"	"0100"	XAXX
	"10"	"0010"	XXAX
	"11"	"0001"	XXAX
MEM_H MEM_HU	"00"	"1100"	BAXX
	"01"	"1100"	BAXX
	"10"	"0011"	XXBA
	"11"	"0011"	XXBA
MEM_W	"00"	"1111"	DCBA
	"01"	"1111"	DCBA
	"10"	"1111"	DCBA
	"11"	"1111"	DCBA

Table 2.10: Computation of M.byteena and M.wrdata, W = DCBA

Operation	A(1 downto 0)	R
MEM_B	"00"	SSSD
	"01"	SSSC
	"10"	SSSB
	"11"	SSSA
MEM_BU	"00"	000D
	"01"	000C
	"10"	000B
	"11"	000A
MEM_H	"00"	SSDC
	"01"	SSDC
	"10"	SSBA
	"11"	SSBA
MEM_HU	"00"	00DC
	"01"	00DC
	"10"	00BA
	"11"	00BA
MEM_W	"00"	DCBA
	"01"	DCBA
	"10"	DCBA
	"11"	DCBA

Table 2.11: Computation of R, D = DCBA

op.memread	op.memtype	A(1 downto 0)	A(ADDR_WIDTH-1 downto 2)	XL
'1'	—	"00"	(others => '0')	'1'
'1'	MEM_H	"01"	—	'1'
'1'	MEM_H	"11"	—	'1'
'1'	MEM_HU	"01"	—	'1'
'1'	MEM_HU	"11"	—	'1'
'1'	MEM_W	"01"	—	'1'
'1'	MEM_W	"10"	—	'1'
'1'	MEM_W	"11"	—	'1'
		otherwise		'0'

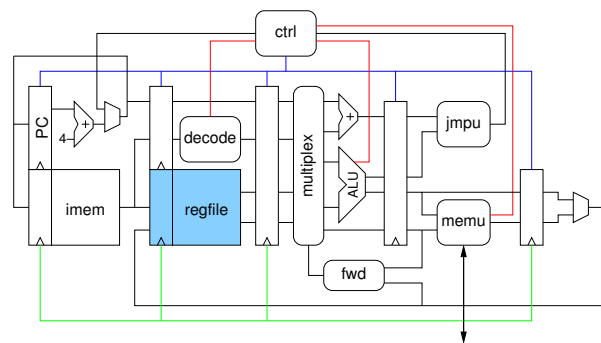
Table 2.12: Memory load exception computation

op.memwrite	op.memtype	A(1 downto 0)	A(ADDR_WIDTH-1 downto 2)	XS
'1'	—	"00"	(others => '0')	'1'
'1'	MEM_H	"01"	—	'1'
'1'	MEM_H	"11"	—	'1'
'1'	MEM_HU	"01"	—	'1'
'1'	MEM_HU	"11"	—	'1'
'1'	MEM_W	"01"	—	'1'
'1'	MEM_W	"10"	—	'1'
'1'	MEM_W	"11"	—	'1'
		otherwise		'0'

Table 2.13: Memory store exception computation

2.4 Register File

regfile.vhd



Description

The register file is a memory with two read ports and one write port, with $2 \times \text{REG_BITS}$ words that are DATA_WIDTH bits wide. The clock signal clk has the usual meaning and causes the circuit to latch the read and write addresses. The reset signal reset is active low and resets internal registers, but not necessarily the contents of the register file. The signal stall causes the circuit not to latch input values such that old values are kept in all registers. Reads from address 0 must always return 0, which may be achieved by an appropriate power-up value and ignoring writes to that location or by intercepting reads from that location. When reading from a register that is written in the same cycle, the new value shall be returned.

In the original MIPS implementation, reads took place on positive clock edges, while writes were performed on negative clock edges in order to forward new values through the register file. However, using both clock edges does not work in the FPGAs used in this lab course. Therefore, the required behavior has to be implemented differently: If the internal register for a read address matches wraddr and $\text{regwrite} = '1'$, the register file shall return wrdata .

Theory

Given memory blocks with one write- and one read-port, how can a memory with one write- and two read-ports be implemented efficiently? What is the overhead, compared to a memory with one write- and one read-port?

Signal	Direction	Type	Width
clk	in	<code>std_logic</code>	—
reset	in	<code>std_logic</code>	—
stall	in	<code>std_logic</code>	—
rdaddr1	in	<code>std_logic_vector</code>	REG_BITS
rdaddr2	in	<code>std_logic_vector</code>	REG_BITS
wraddr	in	<code>std_logic_vector</code>	REG_BITS
wrdata	in	<code>std_logic_vector</code>	DATA_WIDTH
regwrite	in	<code>std_logic</code>	—
rddata1	out	<code>std_logic_vector</code>	DATA_WIDTH
rddata2	out	<code>std_logic_vector</code>	DATA_WIDTH

Table 2.14: Register file interface

3 Level 1: Pipeline

In this assignment, the first version of the pipeline shall be implemented. The pipeline shall be able to execute code, though without resolving any hazards in the pipeline. This means that the results of operations are not available until two cycles later, and that branches have three-cycle branch delay slots.

The pipeline is a classic 5-stage pipeline, consisting of fetch, decode, execute, memory, and write-back stages.

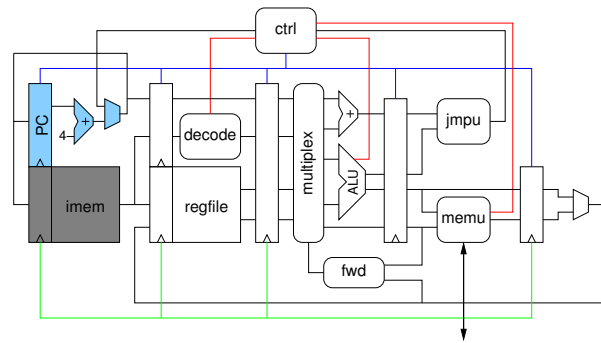
Points 6

Evaluation

The assignment will be tested with test benches, which check the correctness of the behavior at the memory interface for a given content of the instruction memory. Note that this means that testing is only possible if memory operations are implemented. Points will be granted if the design passes the test suites. The assignment is part of lab exercise III.

3.1 Fetch

fetch.vhd



Description

In the fetch stage, the instruction memory is read, and the next value of the program counter is computed. The instruction memory is located within this pipeline stage. Table 3.1 shows the interface of the fetch stage. `clk` and `reset` have their usual meaning, `reset` is active low. After a reset, the fetch stage shall return the instruction located at address 0 in the instruction memory. `stall` causes the fetch stage not to change internal registers, i.e., the program counter must not change while `stall` is asserted. Otherwise, if `pcsrc` is asserted, the next program counter shall be `pc_in`, if `pcsrc` is zero, it shall be the current program counter incremented by 4.

Note that the read port of the instruction memory is registered, which entails that it must be connected to the *next* program counter in order to output the instruction that corresponds to the current program counter register. The *next* program counter is also passed on to the decode stage (see Figure 1.1). Therefore, the program counter in the decode stage does not match the address of the instruction to be decoded, but is usually already incremented by 4. Furthermore, the program counter holds a byte address, while the instruction memory is word-addressed. The lowest two bits of the program counter—which are always zero anyways—are therefore not used to address the instruction memory.

Theory

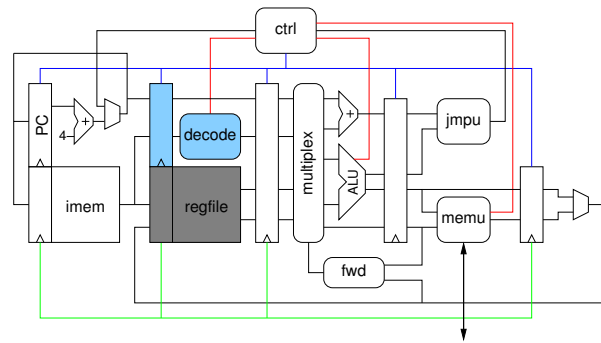
Sketch a fetch stage with variable-length instructions, where the value for the next program counter depends on the instruction that is currently fetched. Which sub-components would be on the critical path in such a fetch stage?

Signal	Dir.	Type	Width	Description
<code>clk</code>	in	<code>std_logic</code>	—	Clock
<code>reset</code>	in	<code>std_logic</code>	—	Reset
<code>stall</code>	in	<code>std_logic</code>	—	Stall
<code>pcsrc</code>	in	<code>std_logic</code>	—	Use <code>pc_in</code> or incremented program counter as new program counter
<code>pc_in</code>	in	<code>std_logic_vector</code>	<code>PC_WIDTH</code>	New program counter
<code>pc_out</code>	out	<code>std_logic_vector</code>	<code>PC_WIDTH</code>	Current program counter
<code>instr</code>	out	<code>std_logic_vector</code>	<code>INSTR_WIDTH</code>	Fetched instruction

Table 3.1: Fetch stage interface

3.2 Decode

decode.vhd



Description

The decode stage contains the register file and translates the raw instructions to signals that are used subsequently in the pipeline. More than one instruction may be mapped to an operation of a function unit such as the ALU. For example, addition of two registers, of a register and an immediate and memory accesses are all mapped to the ALU instruction `ALU_ADD`. Table 3.2 shows the interface of the decode stage. We provide definitions for the types `EXEC_OP_TYPE`, `COP0_OP_TYPE`, `JMP_OP_TYPE`, `MEM_OP_TYPE`, and `WB_OP_TYPE`. You are however free to modify these types in order to optimize your design. The definitions for `JMP_OP_TYPE` and `MEM_OP_TYPE` are provided in Tables 2.6 and 2.8, respectively. `EXEC_OP_TYPE`, `COP0_OP_TYPE` and `WB_OP_TYPE` are described in Tables 3.3, 3.4 and 3.5.

The signals `clk` and `reset` have their usual meaning, `reset` is active low. Asserting `stall` causes the stage not to latch inputs into its internal registers; asserting `flush` causes the unit to store a no-op (all bits cleared) to its internal instruction register.

Figure 3.1 shows the MIPS instruction formats. The operations that the processor must support are shown in Tables 3.6, 3.7, 3.8 and 3.9. The operation semantics in these tables are given in C-syntax. The decoding exception signal `exc_dec` shall be asserted if an instruction cannot be found in one of these tables. As the coprocessor 0 is not implemented at this level, the instructions in Table 3.9 may be treated as no-ops.

Theory

Explain why it is beneficial to have source registers in the same position for all instruction formats, and why this is less of an issue for destination registers.

Signal	Dir.	Type	Width	Description
clk	in	std_logic	–	Clock
reset	in	std_logic	–	Reset
stall	in	std_logic	–	Stall
flush	in	std_logic	–	Flush
pc_in	in	std_logic_vector	PC_WIDTH	Program counter from fetch stage
instr	in	std_logic_vector	INSTR_WIDTH	Instruction to be decoded
wraddr	in	std_logic_vector	REG_BITS	Address for writes to register file
wrdata	in	std_logic_vector	DATA_WIDTH	Data for writes to register file
regwrite	in	std_logic	–	Enable write to register file
pc_out	out	std_logic_vector	PC_WIDTH	Program counter for subsequent stages
exec_op	out	EXEC_OP_TYPE	–	Operation for execute stage
cop0_op	out	COP0_OP_TYPE	–	Operation for coprocessor 0, which handles exceptions and interrupts
jmp_op	out	JMP_OP_TYPE	–	Operation for jump unit
mem_op	out	MEM_OP_TYPE	–	Operation for memory unit
wb_op	out	WB_OP_TYPE	–	Operation for write-back stage
exc_dec	out	std_logic	–	Decoding exception

Table 3.2: Decode stage interface

Field	Type	Width	Description
alu_op	ALU_OP_TYPE	–	ALU operation
readdata1	std_logic_vector	DATA_WIDTH	Data from first register file read port
readdata2	std_logic_vector	DATA_WIDTH	Data from second register file read port
imm	std_logic_vector	DATA_WIDTH	Immediate value from instruction
rs	std_logic_vector	REG_BITS	Value of R-format field rs
rt	std_logic_vector	REG_BITS	Value of R-format field rt
rd	std_logic_vector	REG_BITS	Value of R-format field rd
useimm	std_logic	–	Use immediate value (for ALU or jumps)
useamt	std_logic	–	Use value of shamt field (for shifts only)
link	std_logic	–	Result is (adjusted) value of program counter
branch	std_logic	–	Branch relative to program counter
regdst	std_logic	–	Destination register is in R-format field rt or rd
cop0	std_logic	–	Result is value from coprocessor 0
ovf	std_logic	–	Pass on overflow signal from ALU

Table 3.3: EXEC_OP_TYPE fields

Field	Type	Width	Description
wr	std_logic	–	Write to coprocessor 0 register
addr	std_logic_vector	REG_BITS	Coprocessor 0 register to read from or write to

Table 3.4: COP0_OP_TYPE fields

Field	Type	Description
memtoreg	std_logic	Use ALU or memory result
regwrite	std_logic	Write to register

Table 3.5: WB_OP_TYPE fields

	31	26	25	21	20	16	15	11	10	6	5	0
R-format	opcode	rs		rt		rd		shamt		func		
I-format	opcode	rs		rd		address/immediate						
J-format	opcode	target address										

Figure 3.1: Instruction formats

In Tables 3.6, 3.7, 3.8 and 3.9, apart from C syntax, the following symbols are used:

\emptyset	Unsigned or zero-extended value
\pm	Signed or sign-extended value
$r_{a:b}$	Bits a to b of register r
$[a]$	Value at memory address a

Please also note that in the column labelled “Syntax”, imm18 denotes an 18-bit immediate value with its lowest two bit clear, which enables storing this value in the 16-bit field of the instruction. These values therefore have to be shifted by two bits before being used.

The value pc corresponds to the value of the program counter as it is passed on from the fetch stage, i.e., it corresponds to the next program counter rather than the address of the currently executed instruction.

Opcode	Format	Syntax	Semantics
000000	R	–	see Table 3.7
000001	I	–	see Table 3.8
000010	J	J address	$pc = address^0 \ll 2$
000011	J	JAL address	$r31 = pc+4; pc = address^0 \ll 2$
000100	I	BEQ rd, rs, imm18	if (rs == rd) pc += $imm^{\pm} \ll 2$
000101	I	BNE rd, rs, imm18	if (rs != rd) pc += $imm^{\pm} \ll 2$
000110	I	BLEZ rs, imm18	if ($rs^{\pm} \leq 0$) pc += $imm^{\pm} \ll 2$
000111	I	BGTZ rs, imm18	if ($rs^{\pm} > 0$) pc += $imm^{\pm} \ll 2$
001000	I	ADDI rd, rs, imm16	$rd = rs + imm^{\pm}$, overflow trap
001001	I	ADDIU rd, rs, imm16	$rd = rs + imm^{\pm}$
001010	I	SLTI rd, rs, imm16	$rd = (rs^{\pm} < imm^{\pm}) ? 1 : 0$
001011	I	SLTIU rd, rs, imm16	$rd = (rs^0 < imm^0) ? 1 : 0$
001100	I	ANDI rd, rs, imm16	$rd = rs \& imm^0$
001101	I	ORI rd, rs, imm16	$rd = rs imm^0$
001110	I	XORI rd, rs, imm16	$rd = rs \wedge imm^0$
001111	I	LUI rd, imm16	$rd = imm^0 \ll 16$
010000	R	–	see Table 3.9
100000	I	LB rd, imm16(rs)	$rd = (int8_t)[rs+imm^{\pm}]$
100001	I	LH rd, imm16(rs)	$rd = (int16_t)[rs+imm^{\pm}]$
100011	I	LW rd, imm16(rs)	$rd = (int32_t)[rs+imm^{\pm}]$
100100	I	LBU rd, imm16(rs)	$rd = (uint8_t)[rs+imm^{\pm}]$
100101	I	LHU rd, imm16(rs)	$rd = (uint16_t)[rs+imm^{\pm}]$
101000	I	SB rd, imm16(rs)	$(int8_t)[rs+imm^{\pm}] = rd_{7:0}$
101001	I	SH rd, imm16(rs)	$(int16_t)[rs+imm^{\pm}] = rd_{15:0}$
101011	I	SW rd, imm16(rs)	$(int32_t)[rs+imm^{\pm}] = rd$

Table 3.6: MiMi instructions

Func	Syntax	Semantics
000000	SLL rd, rt, shamt	$rd = rt \ll shamt$
000010	SRL rd, rt, shamt	$rd = rt^0 \gg shamt$
000011	SRA rd, rt, shamt	$rd = rt^{\pm} \gg shamt$
000100	SLLV rd, rt, rs	$rd = rt \ll rs_{4:0}$
000110	SRLV rd, rt, rs	$rd = rt^0 \gg rs_{4:0}$
000111	SRAV rd, rt, rs	$rd = rt^{\pm} \gg rs_{4:0}$
001000	JR rs	$pc = rs$
001001	JALR rd, rs	$rd = pc+4; pc = rs$
100000	ADD rd, rs, rt	$rd = rs + rt$, overflow trap
100001	ADDU rd, rs, rt	$rd = rs + rt$
100010	SUB rd, rs, rt	$rd = rs - rt$, overflow trap
100011	SUBU rd, rs, rt	$rd = rs - rt$
100100	AND rd, rs, rt	$rd = rs \& rt$
100101	OR rd, rs, rt	$rd = rs rt$
100110	XOR rd, rs, rt	$rd = rs \wedge rt$
100111	NOR rd, rs, rt	$rd = \sim(rs rt)$
101010	SLT rd, rs, rt	$rd = (rs^{\pm} < rt^{\pm}) ? 1 : 0$
101011	SLTU rd, rs, rt	$rd = (rs^0 < rt^0) ? 1 : 0$

Table 3.7: MiMi special instructions

rd	Syntax	Semantics
00000	BLTZ rs, imm18	if $(rs^{\pm} < 0) pc += imm^{\pm} \ll 2$
00001	BGEZ rs, imm18	if $(rs^{\pm} \geq 0) pc += imm^{\pm} \ll 2$
10000	BLTZAL rs, imm18	$r31 = pc+4$; if $(rs^{\pm} < 0) pc += imm^{\pm} \ll 2$
10001	BGEZAL rs, imm18	$r31 = pc+4$; if $(rs^{\pm} \geq 0) pc += imm^{\pm} \ll 2$

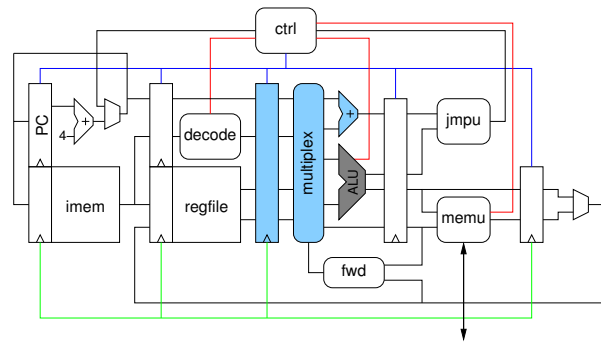
Table 3.8: MiMi regimm instructions

rs	Syntax	Semantics
00000	MFC0 rt, rd	$rt = rd$, rd register in coprocessor 0
00100	MTC0 rt, rd	$rd = rt$, rd register in coprocessor 0

Table 3.9: MiMi cop0 instructions

3.3 Execute

exec.vhd



Description

The execute stage contains the ALU, and therefore “executes” the arithmetic and logic instructions. Furthermore, the ALU is used to compute the addresses for memory accesses. Also, the addition for branches relative to the program counter is computed in this stage. Table 3.10 shows the interface of the execute stage.

The signals `clk` and `reset` have their usual meaning, `reset` is active low. Asserting `stall` causes the stage not to latch inputs into its internal registers; asserting `flush` causes the unit to store no-ops to the pipeline registers. The signal `exc_ovf` shall be asserted if the ALU asserts the overflow flag `V` and the current operation may trigger an overflow trap.

The signals `rs` and `rt` shall be instruction fields `rs` and `rt`, respectively. The signal `rd` shall be the destination register for the current operation, which may correspond to the field `rd` or `rt`, depending on the instruction format.

For most instructions, the `aluresult` signal is the result from the ALU. For the `mfc0` instruction, it shall hold the result from coprocessor 0. For instructions such as `jal` or `jalr`, `aluresult` shall contain the (adjusted) program counter. Please note that for the `bltzal` and `bgtzal` instructions, it is not possible to use the ALU to both compute the condition and adjust the program counter. Either the zero and neg flags have to be computed separately, or a separate adder to adjust the program counter has to be used.

Information in the signals suffixed `_in` and `_out` shall be passed on to subsequent pipeline stages without being modified.

The signals `forwardA`, `forwardB`, `mem_aluresult` and `wb_result` are irrelevant for this assignment and can be ignored. They will be used for forwarding the correct data to the ALU in lab exercise IV.

Theory

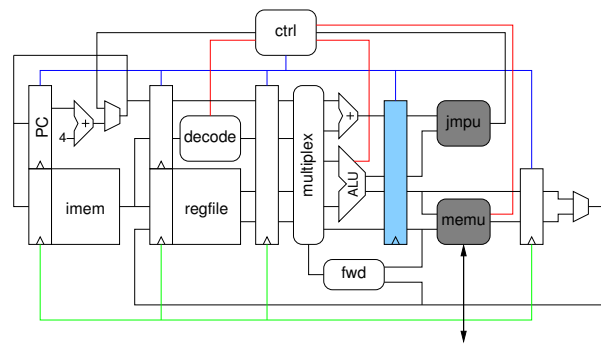
Explain why it is beneficial to multiplex the operands for a single adder over using several adders and multiplex their results. Does the benefit concern rather the performance or the size of the resulting hardware?

Signal	Dir.	Type	Width	Description
clk	in	std_logic	—	Clock
reset	in	std_logic	—	Reset
stall	in	std_logic	—	Stall
flush	in	std_logic	—	Flush
op	in	EXEC_OP_TYPE	—	Operation for this stage
rd	out	std_logic_vector	REG_BITS	Value of instruction's destination field
rs	out	std_logic_vector	REG_BITS	Value of instruction's rs field
rt	out	std_logic_vector	REG_BITS	Value of instruction's rt field
alurestult	out	std_logic_vector	DATA_WIDTH	Result from ALU or coprocessor 0, or adjusted PC
wrdata	out	std_logic_vector	DATA_WIDTH	Value to be written to memory
zero	out	std_logic	—	Zero flag from ALU
neg	out	std_logic	—	Negative result from ALU
new_pc	out	std_logic_vector	PC_WIDTH	Target address for branches
pc_in	in	std_logic_vector	PC_WIDTH	Program counter from decode stage
pc_out	out	std_logic_vector	PC_WIDTH	Program counter to memory stage
memop_in	in	MEM_OP_TYPE	—	Memory operation from decode stage
memop_out	out	MEM_OP_TYPE	—	Memory operation to memory stage
jmpop_in	in	JMP_OP_TYPE	—	Jump operation from decode stage
jmpop_out	out	JMP_OP_TYPE	—	Jump operation to memory stage
wbop_in	in	WB_OP_TYPE	—	Write-back operation from decode stage
wbop_out	out	WB_OP_TYPE	—	Write-back operation to memory stage
forwardA	in	FWD_TYPE	—	Forwarding info for operand A
forwardB	in	FWD_TYPE	—	Forwarding info for operand B
cop0_rddata	in	std_logic_vector	DATA_WIDTH	Data from coprocessor 0
mem_alurestult	in	std_logic_vector	DATA_WIDTH	Result from ALU from previous cycle, from memory stage
wb_result	in	std_logic_vector	DATA_WIDTH	Result from ALU two cycles ago or from memory operation, from write-back stage
exc_ovf	out	std_logic	—	Overflow exception

Table 3.10: Execute stage interface

3.4 Memory

mem.vhd



Description

Despite its name, the memory stage does not only contain the memory unit, but also the jump unit. Most of its functionality is already implemented in these two units. Therefore, the implementation for this stage mainly consists of registering the inputs and passing them on to the memory and jump unit. The interface for this stage is shown in Table 3.11.

The signals `clk` and `reset` have their usual meaning, `reset` is active low. Asserting `flush` causes the unit to store no-ops to the pipeline registers. Asserting `stall` causes the stage not to latch inputs into its internal registers; additionally, neither `op.memread` nor `op.memwrite` of the memory unit may be asserted while the `stall` signal is asserted.

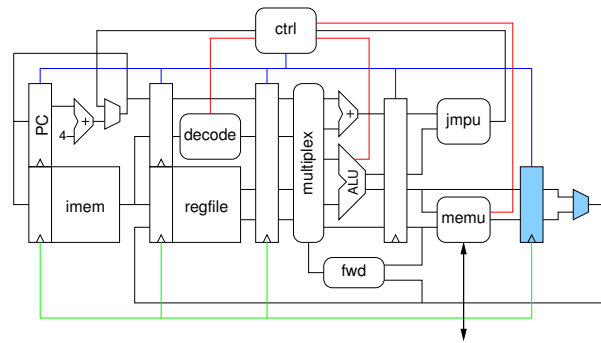
Information in the signals suffixed `_in` and `_out` shall be passed on to subsequent pipeline stages without being modified. Other signals shall be connected to the appropriate ports of the jump and memory units.

Signal	Dir.	Type	Width	Description
clk	in	std_logic	–	Clock
reset	in	std_logic	–	Reset
stall	in	std_logic	–	Stall
flush	in	std_logic	–	Flush
mem_op	in	MEM_OP_TYPE	–	Memory operation from execute stage
jmp_op	in	JMP_OP_TYPE	–	Jump operation from execute stage
wrdata	in	std_logic_vector	DATA_WIDTH	Data to be written to memory
memresult	out	std_logic_vector	DATA_WIDTH	Result of memory load
zero	in	std_logic	–	Zero flag from ALU
neg	in	std_logic	–	Negative result from ALU
pcsrc	out	std_logic	–	Asserted if a jump is to be executed
new_pc_in	in	std_logic_vector	PC_WIDTH	Jump target from execute stage
new_pc_out	out	std_logic_vector	PC_WIDTH	Jump target to fetch stage
pc_in	in	std_logic_vector	PC_WIDTH	Program counter from execute stage
pc_out	out	std_logic_vector	PC_WIDTH	Program counter to write-back stage
rd_in	in	std_logic_vector	REG_BITS	Destination register from execute stage
rd_out	out	std_logic_vector	REG_BITS	Destination register to write-back stage
aluresult_in	in	std_logic_vector	DATA_WIDTH	Result from ALU from execute stage
aluresult_out	out	std_logic_vector	DATA_WIDTH	Result from ALU to write-back stage
wbop_in	in	WB_OP_TYPE	–	Write-back operation from execute stage
wbop_out	out	WB_OP_TYPE	–	Write-back operation to write-back stage
mem_out	out	MEM_OUT_TYPE	–	Memory operation to outside the pipeline
mem_data	in	std_logic_vector	DATA_WIDTH	Memory load result from outside the pipeline
exc_load	out	std_logic	–	Load exception
exc_store	out	std_logic	–	Store exception

Table 3.11: Memory stage interface

3.5 Write-Back

wb.vhd



Description

The purpose of the write-back stage is to select between the result from the ALU and from memory loads and to relax the critical paths in the pipeline. Table 3.12 shows its interface.

Signal	Dir.	Type	Width	Description
clk	in	std_logic	—	Clock
reset	in	std_logic	—	Reset
stall	in	std_logic	—	Stall
flush	in	std_logic	—	Flush
op	in	WB_OP_TYPE	—	Write-back operation from memory stage
alurestult	in	std_logic_vector	DATA_WIDTH	Result from ALU
memresult	in	std_logic_vector	DATA_WIDTH	Result from memory load
result	out	std_logic_vector	DATA_WIDTH	Result to register file
regwrite	out	std_logic	—	Write enable to register file
rd_in	in	std_logic_vector	REG_BITS	Destination register from memory stage
rd_out	out	std_logic_vector	REG_BITS	Destination register to register file

Table 3.12: Write-back stage interface

3.6 Pipeline

pipeline.vhd

Description

The pipeline stages described above shall be connected to form a pipeline. The interface of the pipeline is shown in Table 3.13. The `clk` and `reset` signals have their usual meaning, `reset` is active low. If the field `busy` in the input signal `mem_in` is asserted, the pipeline shall be stalled. As the `ctrl` unit is not yet implemented, the appropriate signals from the memory stage shall be passed on to the fetch stage without modifying them. As the pipeline in its current state does not resolve *any* hazards, the `flush` signal of the individual pipeline stages can be hardwired to '0'. The signal `intr` can be ignored for this assignment, but will be used in lab exercise IV to trigger external interrupts.

Signal	Dir.	Type	Width	Description
<code>clk</code>	in	<code>std_logic</code>	–	Clock
<code>reset</code>	in	<code>std_logic</code>	–	Reset
<code>mem_in</code>	in	<code>MEM_IN_TYPE</code>	–	Interface from memory to the pipeline
<code>mem_out</code>	out	<code>MEM_OUT_TYPE</code>	–	Interface from the pipeline to the memory
<code>intr</code>	in	<code>std_logic_vector</code>	<code>INTR_COUNT</code>	External interrupt lines

Table 3.13: Pipeline interface

The pipeline should now be able to execute sequences of assembly code. As hazards are not resolved, the results from operations only become available two instructions later. Also, branches require a three-cycle branch delay slot. The assembler code shown in Listing 1 shows an endless loop that stores the numbers 0, 1, 2, ... to address 16. Note that after initializing or incrementing register `$1` two `nop` operations are necessary for correct operation.

```

1      addi $1, $0, 0
2      nop
3      nop
4 loop:
5      addi $1, $1, 1
6      nop
7      nop
8      sw $1, 16($0)
9      j loop
10     nop
11     nop
12     nop

```

Listing 1: Assembler example without forwarding

Theory

Listing 1 contains seven `nop`-instructions. How many of these instructions can be removed by reordering instructions, without changing the semantics of the program?

A Tools

Listing 2 shows a basic Makefile to compile programs for MiMi in the lab environment. The variable `CC` denotes the C compiler, `LD` the linker, `AR` the program to create library files and `OBJCOPY` the program to convert between binary representations of a program. The default flags for the compiler indicate that MiMi is compatible to the MIPS 1 ISA and that it does not contain a floating-point unit. The flags for `LD` determine the memory layout and must not be changed.

The target `lib` builds the program preamble `crt0.o` and a minimal version of the C library, `libc.a`. This target must be built explicitly before building programs. Note that it is necessary to provide an explicit rule for linking in order to link the right files (the `libc.a` in the current directory) in the right order (`crt0.o` before everything else).

The build process for MiMi is slightly more complex than for more common architectures. After having linked the program in the ELF file format, the binary needs further processing. First, the data for the instruction and data memories have to be extracted to files which end in `.imem.hex` and `.dmem.hex`, respectively. Then, these files have to be converted from the Intel HEX format to the MIF file format understood by Quartus, which is done with the `hex2mif.pl` script. The `.mif` files then need to be copied to the appropriate directory such that the correct program is synthesized into the processor.

```
1 PREFIX=/usr/mips
2
3 CC=${PREFIX}/bin/mips-elf-gcc -mips1 -msoft-float
4 LD=${PREFIX}/bin/mips-elf-ld -N -Ttext=0x40000000 --section-start .rodata=4
5 AR=${PREFIX}/bin/mips-elf-ar
6 OBJCOPY=${PREFIX}/bin/mips-elf-objcopy
7
8 CFLAGS=-O2 -DARCH_IS_BIG_ENDIAN=1
9
10 test.elf: test.o
11     ${LD} -o $@ crt0.o $^ -L. -lc
12
13 lib: crt0.o libc.a
14
15 libc.a: exceptions.o util.o
16     ${AR} rc $@ $^
17
18 %.imem.hex : %.elf
19     ${OBJCOPY} -j .text -O ihex $< $@
20
21 %.dmem.hex : %.elf
22     ${OBJCOPY} -R .text -O ihex $< $@
23
24 %.mif : %.hex
25     ./hex2mif.pl < $< > $@
```

Listing 2: Makefile fragment

B Automated Test Environment

In order to submit your source code to the automated test system, copy your design to the directory `/ddcanightly/ddcagrp<grpnr>/level<lvlnr>/`. The VHDL files should reside in the subdirectory `src`. For example, group 99 should copy their implementation of the decode stage to `/ddcanightly/ddcagrp99/level1/src/decode.vhd` to submit it to the level 1 tests.

The source code for each level must be complete; higher levels must include the source code from lower levels. The files `imem_altera.vhd`, `ocram_altera.vhd` and `serial_port_wrapper.vhd` must be identical to the provided files. Only the files that were provided to you are taken into account by the test suite. Your design will therefore not pass the test benches if it requires any additional source files.

The automated test benches create a snapshot of the `/ddcanightly` directory daily at 1:00 AM; please do not submit files around that time to avoid an inconsistent state of your source code. The test bench results are reported to the e-mail addresses stated in `/ddcanightly/ddcagrp<grpnr>/recipients`. You will receive the reports only if you enter your e-mail address in that file.

The test system is not intended to replace your own testing. In order to avoid abuse of the test system for debugging, the test reports provide only minimal information on the failed test cases. For test cases that are based on assembler or C code, you will however be provided with the source code.

In order to receive points for the exercises, the source code must be uploaded to the TUWEL system before the deadline. Points will only be awarded if the source code submitted to the TUWEL system passes the test suites.

C Submission Requirements

C.1 Exercise III

The results have to be submitted via TUWEL (Deadline: 25.05.2018, 23:55). Upload a **tar.gz** archive named **submission_ex3.tar.gz** containing the following items:

- Your lab protocol as PDF file.
- The complete VHDL source code for the assignments detailed in Section 2 and 3, including entities and packages provided to you.
Source code must extract to a directory named **src**. It is permissible to place the source code for Level 0 in a directory **level0/src** and the source code for Level 1 in a directory **level1/src**.
- Create an environment to automatically simulate and test the ALU. The simulation should be text-based and does not have to plot any waveform. As a template, the test environment for the ps2_ascii component of exercise I can be used.
By executing "make compile", the ALU VHDL source should be compiled. Afterwards, by executing "make sim_alu" the corresponding testbench should be compiled and the simulation should be performed checking the output of the ALU for the input being provided. The inputs and expected outputs should be read from a file.

The submitted archive should have the following structure:

```

submission_ex3.tar.gz
├── report.pdf..... Include your lab report here
├── src.....The complete VHDL source code.
│   ├── alu.vhd
│   ├── core.vhd
│   └── [... all other source files]
├── test_ALU.....The testing environment for the ALU.
│   ├── scripts
│   │   ├── compile.do
│   │   └── sim_alu.do
│   ├── tb
│   │   └── alu_tb_fileio.vhd
│   ├── testdata
│   │   ├── input.txt
│   │   └── output.txt
│   └── Makefile

```

Acknowledgements

This document was written by Wolfgang Puffitsch. Other people, who have helped in improving it: Jomy Joseph Chelackal, Florian Ferdinand Huemer, Thomas Preindl, Jörg Rohringer, Markus Schütz and others.

References

- [1] John L. Hennessy and David A. Patterson. *Computer Architecture, Fourth Edition: A Quantitative Approach*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2006.
- [2] David A. Patterson and John L. Hennessy. *Computer Organization and Design, Fourth Edition, Fourth Edition: The Hardware/Software Interface (The Morgan Kaufmann Series in Computer Architecture and Design)*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 4th edition, 2008.